


ARTIFICIAL INTELLIGENCE

BY

ENG.Kholoud Amer

THE CONTENT

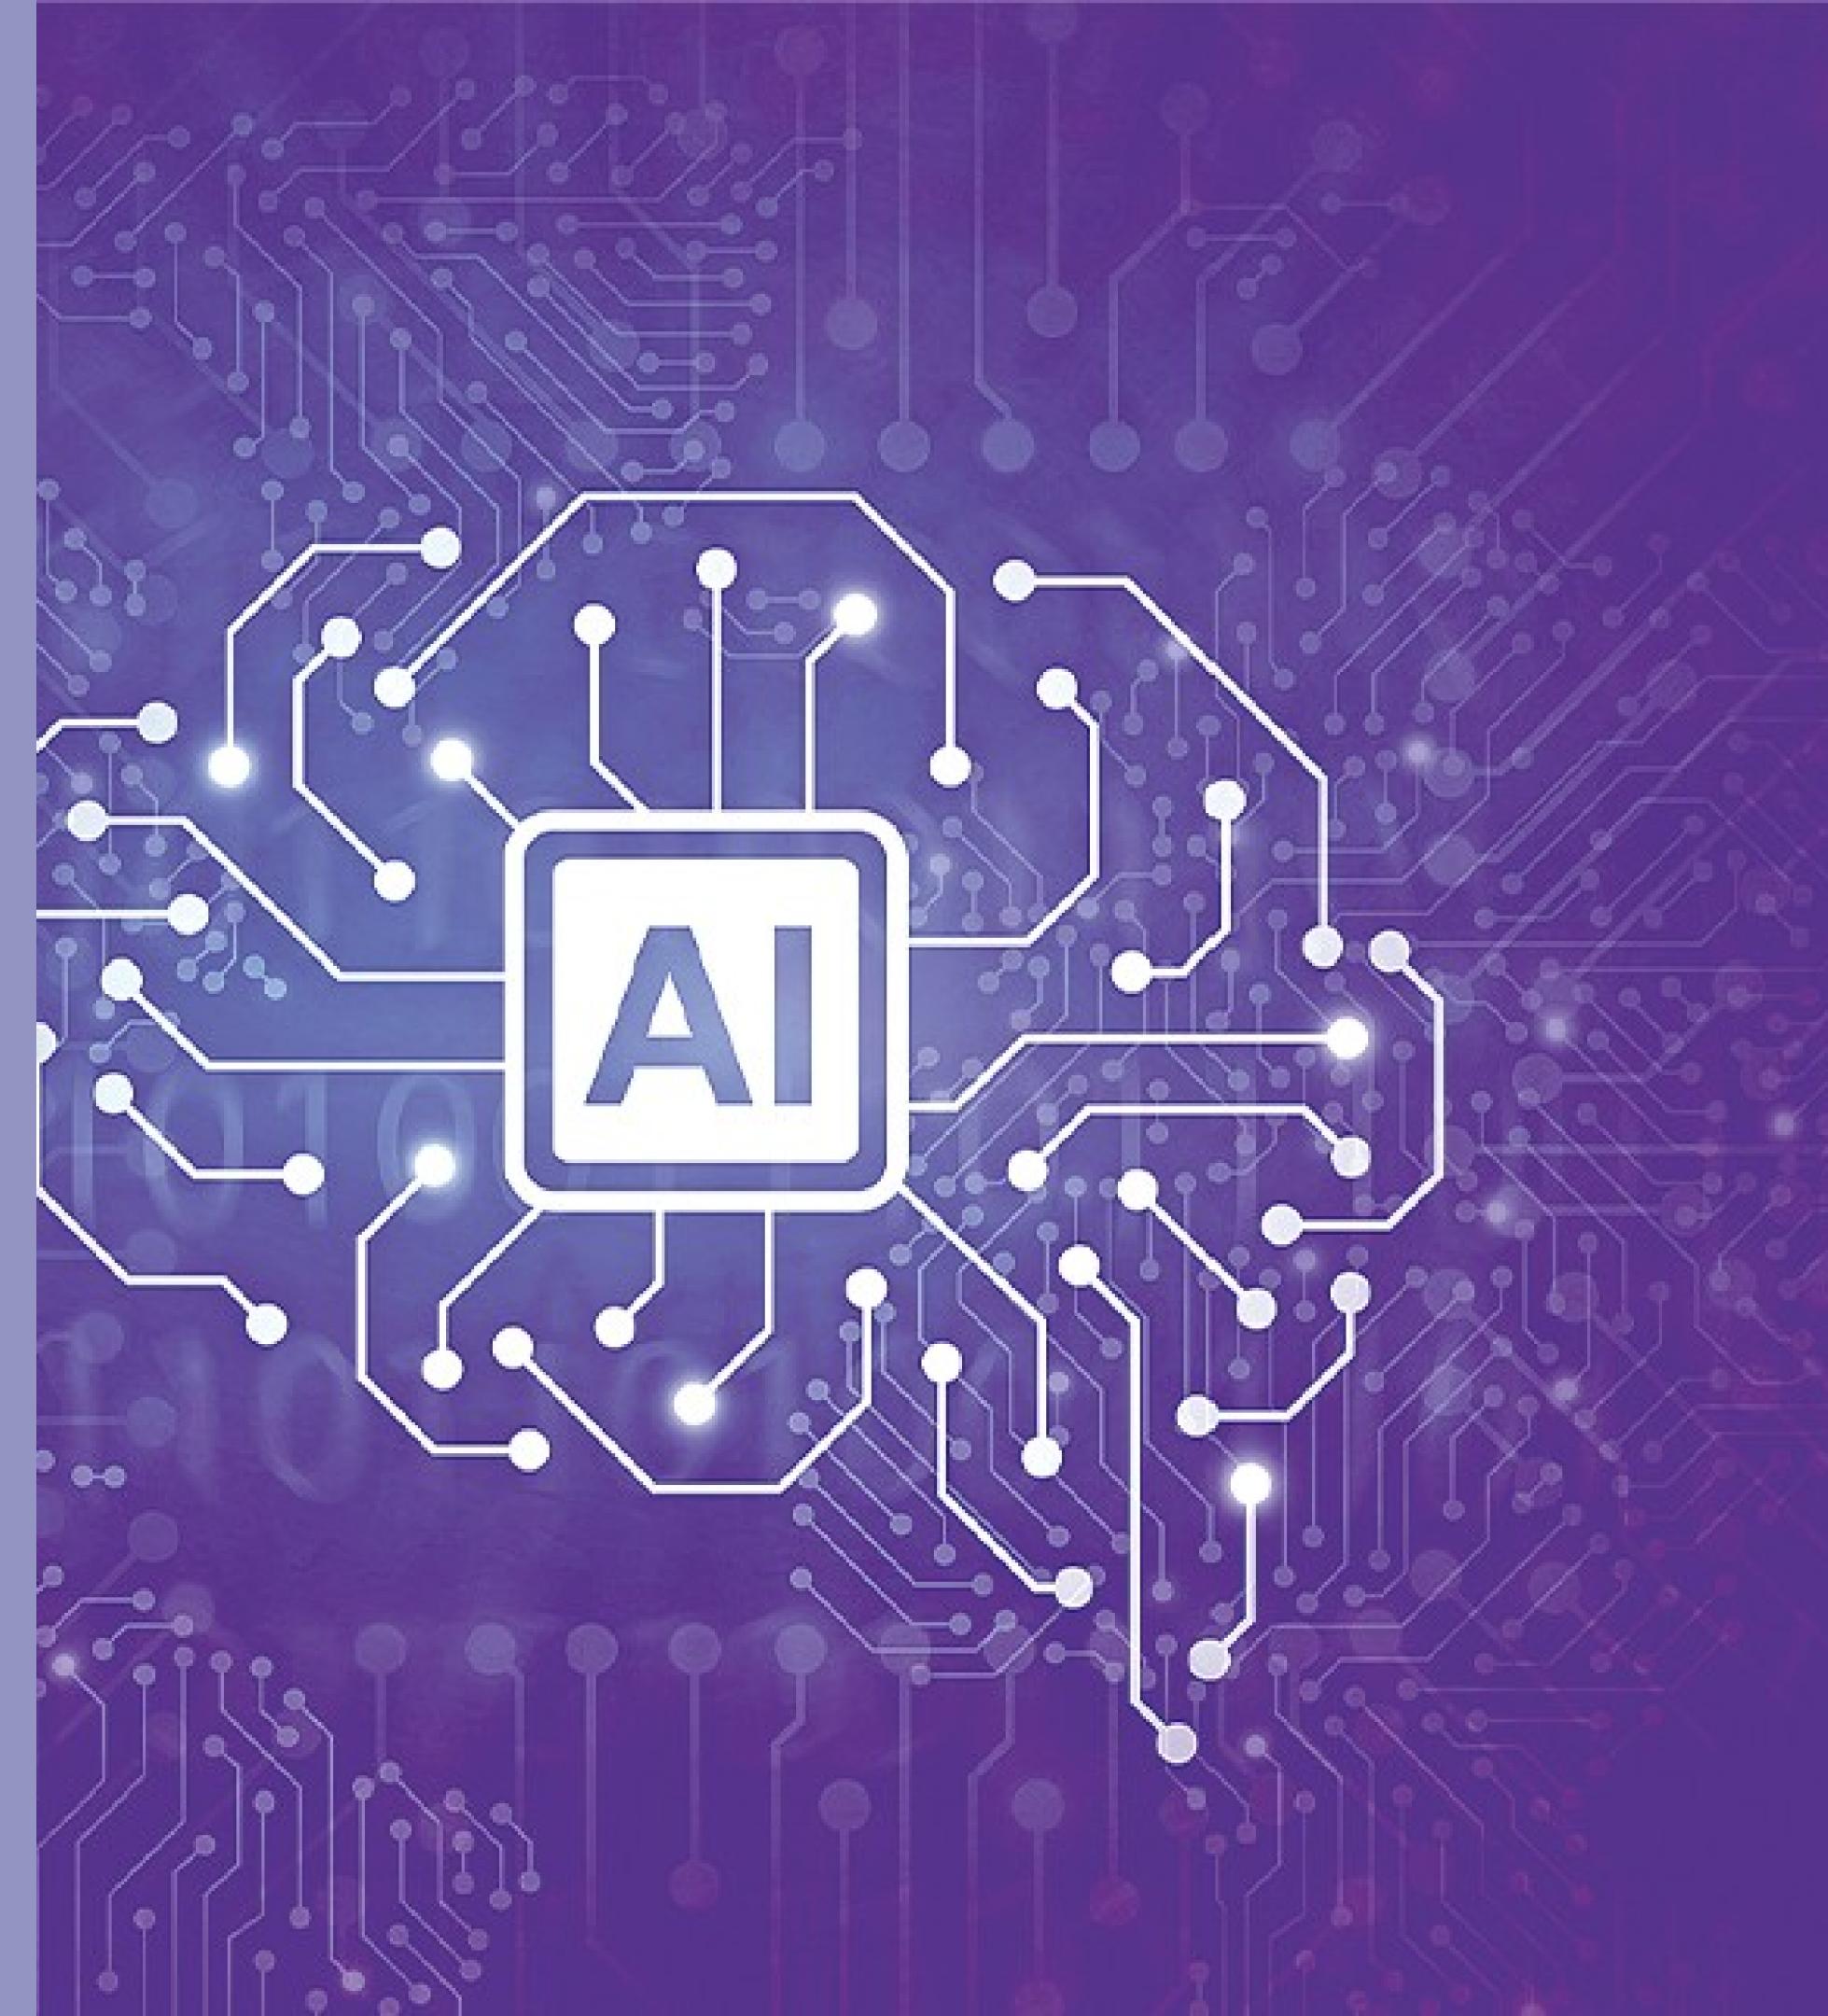


ARTIFICIAL INTELLIGENCE

- 1** What is Deep Learning?
- 2** Artificial Neural Networks
- 3** Convolutional Neural Networks (CNNs)

WHAT IS DEEP LEARNING?

- Deep learning is a subfield of artificial intelligence (AI) that focuses on training computer systems to learn and make decisions in a manner similar to the human brain.
- It is based on artificial neural networks, which are inspired by the structure and function of biological neural networks found in the human brain.
- Deep learning models are capable of automatically learning to represent and understand complex patterns and features from large amounts of data.



WHAT IS DEEP LEARNING?

- Deep learning is a branch of machine learning that employs artificial neural networks with multiple layers (hence "deep") to model and process data.
- These networks are designed to automatically learn and improve from experience, enabling them to perform tasks like image and speech recognition, natural language processing, autonomous driving, and more.

IMPORTANCE AND APPLICATIONS OF DEEP LEARNING

- Deep learning has revolutionized various industries due to its ability to extract meaningful information from massive datasets. Some key applications of deep learning include:**

1. Image and Object Recognition: Deep learning powers applications like facial recognition, image classification, and object detection, leading to advancements in fields like computer vision.


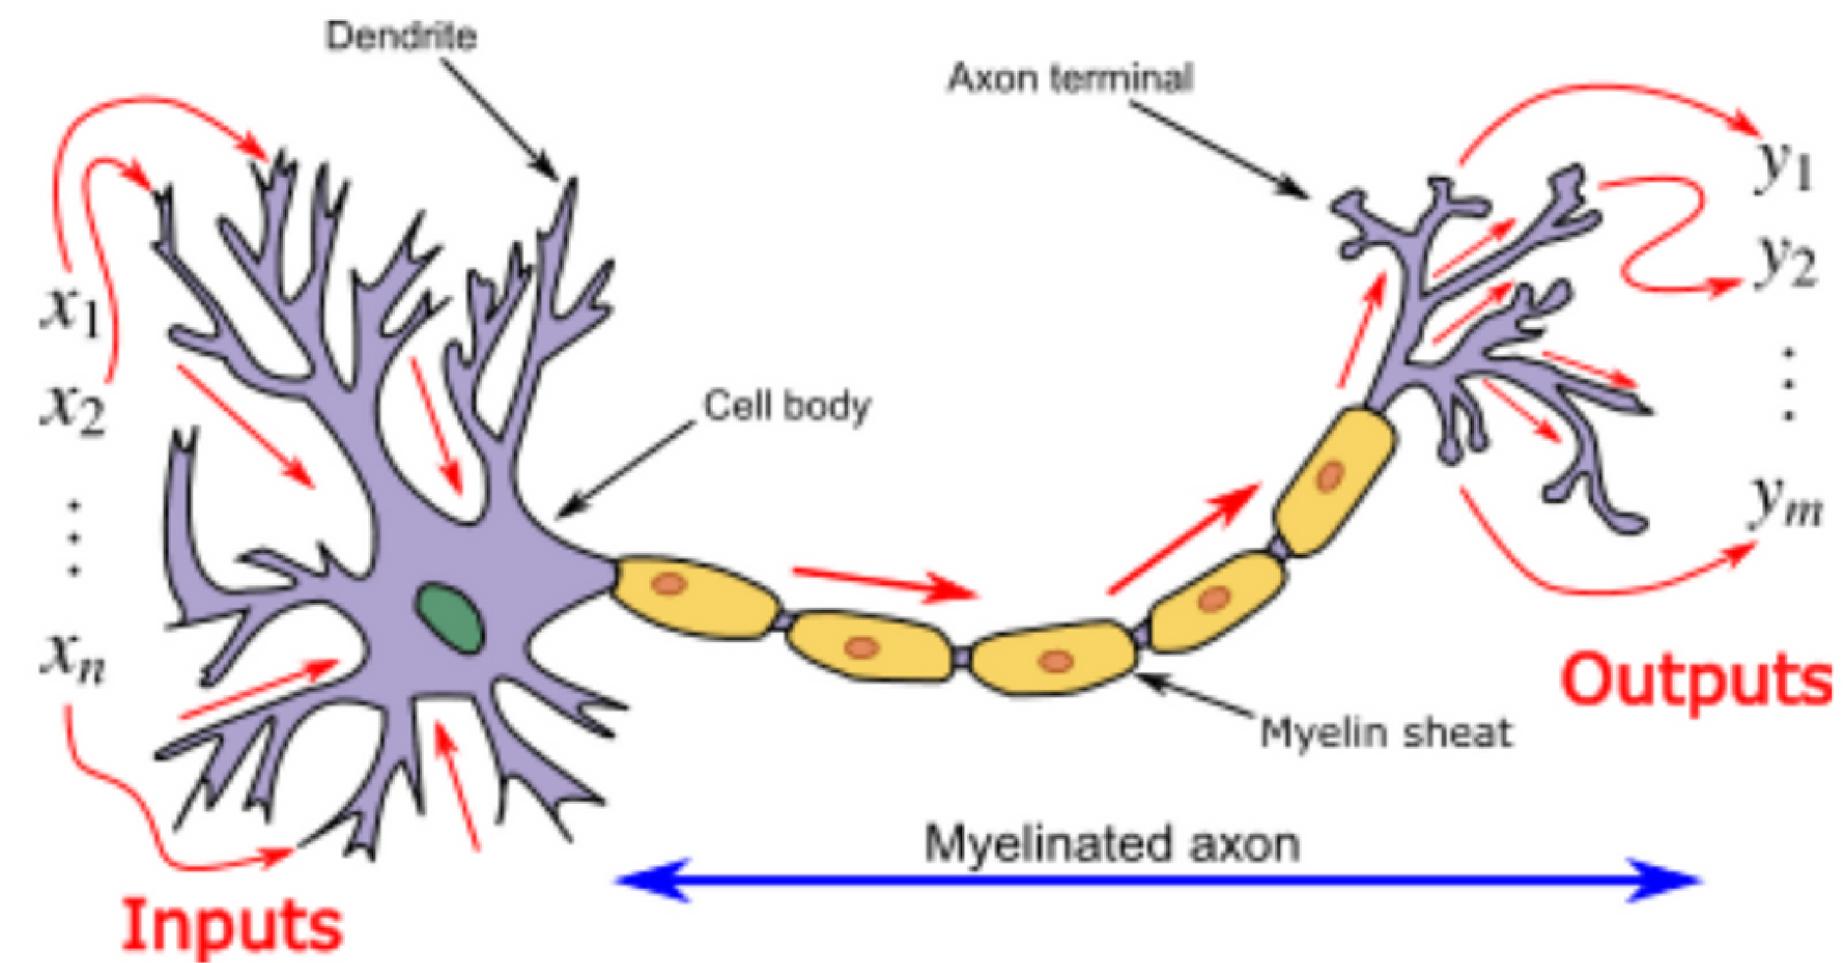
IMPORTANCE AND APPLICATIONS OF DEEP LEARNING

2. Image and Object Recognition: Deep learning powers applications like facial recognition, image classification, and object detection, leading to advancements in fields like computer vision.

3. Natural Language Processing (NLP): NLP enables machines to understand, interpret, and generate human language, enabling chatbots, language translation, sentiment analysis, and more.

IMPORTANCE AND APPLICATIONS OF DEEP LEARNING

4. Autonomous Systems: Deep learning is at the core of self-driving cars and other autonomous machines that can perceive and navigate the world.
5. Autonomous Systems: Deep learning is at the core of self-driving cars and other autonomous machines that can perceive and navigate the world.
6. Healthcare: Deep learning assists in medical image analysis, disease diagnosis, drug discovery, and personalized treatment planning.



ARTIFICIAL NEURAL NETWORKS (ANNs)

- Artificial Neural Networks (ANNs) are computational models inspired by the structure and functioning of the human brain. They consist of interconnected nodes, known as neurons, organized in layers.
- In biology, a neuron is a specialized cell found in the nervous system of animals, including humans. Neurons are responsible for transmitting information through electrical and chemical signals, enabling the brain and nervous system to process and respond to various stimuli.
- Each biological neuron consists of a cell body, dendrites (which receive signals from other neurons), an axon (which transmits signals to other neurons), and synapses (connections between neurons where information is transmitted through neurotransmitters).

ARTIFICIAL NEURAL NETWORKS (ANNs)

Artificial Neuron:

In the context of artificial intelligence, an **artificial neuron** (also known as a "perceptron" or "node") is a simplified mathematical model inspired by biological neurons. It is the fundamental building block of **artificial neural networks**, which are computational models used for various machine learning tasks.

An **artificial neuron** takes multiple input values, multiplies each input by a corresponding weight, and then sums up these weighted inputs. The weighted sum is passed through an activation function to produce an output. The activation function introduces non-linearity to the model, enabling the **artificial neuron** to learn complex patterns in the data.

ARTIFICIAL NEURAL NETWORKS (ANNs)

The layers in an artificial neural network are typically classified into three types:

- 1. Input Layer: The first layer of the network, where data is fed into the network.**
- 2. Hidden Layers: Intermediate layers between the input and output layers. They are responsible for learning and representing complex patterns in the data.**
- 3. Output Layer: The final layer of the network, which produces the network's predictions or outputs based on the learned patterns from the hidden layers.**

ARTIFICIAL NEURAL NETWORKS (ANNs)

Training an artificial neural network involves feeding it labeled data (input with corresponding output) and adjusting the weights of the connections iteratively to minimize the difference between the predicted output and the actual output. This process is often carried out using optimization techniques like backpropagation and gradient descent.

DEEP NEURAL NETWORKS (DNNs)

- Deep Neural Networks (DNNs) are a special type of Artificial Neural Network (ANN) that consist of multiple layers of interconnected artificial neurons.
- These networks are inspired by the structure and functioning of the human brain, and their deep architecture allows them to learn and represent complex patterns in data.

DEEP NEURAL NETWORKS (DNNs)

Going Deeper - Layers and Neurons:

- A DNN typically consists of an input layer, one or more hidden layers, and an output layer.
- Each layer is made up of artificial neurons, also known as nodes, which process and pass information to the next layer.

HOW DNNs LEARN

- DNNs learn by adjusting the connection strengths between neurons, called weights, during a training process.
- The training involves showing the network lots of examples with correct answers and updating the weights to minimize prediction errors.
- Through this iterative process, the DNN becomes better at making accurate predictions on its own.

CONVOLUTIONAL NEURAL NETWORKS (CNNs)

- **Convolutional Neural Networks (CNNs) are a specialized type of Deep Neural Network designed for processing and analyzing visual data, such as images and videos.**
- **They are inspired by how the human visual system works and have revolutionized computer vision tasks.**

KEY COMPONENTS OF CNNS

- **CNNs have three main components: convolutional layers, pooling layers, and fully connected layers.**
- **Convolutional layers use filters (also known as kernels) to scan the input image, detecting patterns and features.**
- **Pooling layers reduce the spatial dimensions of the feature maps, reducing the number of parameters and improving computational efficiency.**
- **Fully connected layers process the extracted features to make final predictions.**

UNDERSTANDING CONVOLUTION

- Convolution is a mathematical operation used in CNNs to extract patterns and features from images.
- Filters are small windows that slide over the image, and at each location, they perform element-wise multiplication and summation to produce a new feature map.
- This process allows the network to learn important patterns, such as edges, corners, and textures, at different scales.
-

FEATURE MAPS AND ACTIVATION MAPS

- Feature maps are the result of applying convolutional filters to the input image.
- Activation maps are obtained by applying an activation function to the feature maps, introducing non-linearity and making the network capable of learning complex relationships.

TRANSFER LEARNING

- Transfer learning is a technique where a pre-trained CNN model is used as a starting point for a new task.
- By leveraging knowledge from previously learned features, transfer learning enables faster and more accurate training on new datasets, especially when data is limited

```
import tensorflow as tf
from tensorflow.keras import layers, models

# Load the MNIST dataset
mnist = tf.keras.datasets.mnist
(train_images, train_labels), (test_images, test_labels)
```

MNIST.LOAD_DATAO

```
# Normalize the pixel values to be between 0 and 1  
train_images, test_images = train_images / 255.0, test_images / 255.0
```

```
# Reshape the images to add the channel dimension (since they are grayscale)  
train_images = train_images.reshape(train_images.shape[0], 28, 28, 1)  
test_images = test_images.reshape(test_images.shape[0], 28, 28, 1)
```

```
# Create the CNN model  
model = models.Sequential()
```

```
# Normalize the pixel values to be between 0 and 1
train_images, test_images = train_images / 255.0, test_images / 255.0

# Reshape the images to add the channel dimension (since they are grayscale)
train_images = train_images.reshape(train_images.shape[0], 28, 28, 1)
test_images = test_images.reshape(test_images.shape[0], 28, 28, 1)

# Create the CNN model
model = models.Sequential()

# Convolutional layer 1
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
```

```
# Convolutional layer 2  
model.add(layers.Conv2D(64, (3, 3), activation='relu'))  
model.add(layers.MaxPooling2D((2, 2)))  
  
# Flatten the output before the fully connected layers  
model.add(layers.Flatten())
```

```
# Fully connected layers  
model.add(layers.Dense(64, activation='relu'))  
model.add(layers.Dense(10, activation='softmax')) # C  
  
# Compile the model  
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```

```
# Train the model  
model.fit(train_images, train_labels, epochs=5, batch_size=32, validation_data=  
          (test_images, test_labels))
```

```
# Evaluate the model on the test set  
test_loss, test_accuracy = model.evaluate(test_images, test_labels)  
print("Test accuracy:", test_accuracy)
```

INTRODUCTION TO RECURRENT NEURAL NETWORKS (RNNs)

In traditional neural networks, data is treated as independent and unrelated. But what if we encounter data with a sequence or temporal dependency?

INTRODUCTION TO RECURRENT NEURAL NETWORKS (RNNs)

- **Sequential data refers to data that has a natural order and is dependent on previous elements.**
- **Examples of sequential data include time series data, speech, text, music, and more.**
- **RNNs are designed to work with such sequential data and excel in capturing patterns across time.**

RECURRENT NEURAL NETWORKS (RNNs) KEY CONCEPTS

- [The Recurrent Neuron]
- A fundamental unit in an RNN is the recurrent neuron, which processes both the current input and the output from the previous time step.
- This connection between previous and current states enables RNNs to maintain a memory of past information, making them powerful for sequential tasks.
-

RECURRENT NEURAL NETWORKS (RNNs) KEY CONCEPTS

- **[Unrolling the RNN]**
- **To process a sequence, an RNN is unrolled through time, creating a chain of recurrent neurons.**
- **Each time step is linked to the previous one, and information flows through the network from one step to another.**
-

RECURRENT NEURAL NETWORKS (RNNs) KEY CONCEPTS

- By unrolling the RNN, the computer can process the sequential data more effectively and learn from the patterns and relationships between different time steps.
- This makes the RNN great for tasks like predicting what might come next in a sequence, understanding the meaning of sentences, or generating new sequences, like music or text
- This way, the RNN can work its magic and do incredible things with sequential data, just like you can enjoy and understand the entire story on that long, unrolled scroll!

RECURRENT NEURAL NETWORKS (RNNs) KEY CONCEPTS

- By unrolling the RNN, the computer can process the sequential data more effectively and learn from the patterns and relationships between different time steps.
- This makes the RNN great for tasks like predicting what might come next in a sequence, understanding the meaning of sentences, or generating new sequences, like music or text
- This way, the RNN can work its magic and do incredible things with sequential data, just like you can enjoy and understand the entire story on that long, unrolled scroll!

HOW THE RNN ALGORITHM WORKS STEP BY STEP

Step 1: Input Data

- The RNN algorithm starts with a sequence of input data, such as a sentence or a time series of values (e.g., stock prices over time).
- Each element in the sequence is called a "time step."

Step 2: Initialization

- Before processing the sequence, the RNN initializes its internal memory to a starting state, usually with all values set to zero or random small numbers.

HOW THE RNN ALGORITHM WORKS STEP BY STEP

Step 3: Unrolling the RNN

- The RNN unrolls over time, just like you unroll a scroll to see each part of the story.
- At each time step, the RNN takes in the input data for that step and combines it with the internal memory from the previous time step.

Step 4: Processing the Data

- The RNN processes the input data at each time step using a special neuron called a "recurrent neuron."
- The recurrent neuron takes both the current input and the previous time step's memory as input and performs a mathematical operation to produce an output.

HOW THE RNN ALGORITHM WORKS STEP BY STEP

Step 5: Updating Memory

- The output of the recurrent neuron becomes the new memory for the current time step.
- This way, the RNN carries information from one time step to the next, just like remembering parts of the story as you read it.

Step 6: Repeating the Process

- The RNN repeats this process for all the time steps in the sequence, gradually updating its memory and producing outputs at each step.

Step 7: Final Output

- After processing all the time steps, the RNN produces an output for each time step, representing its understanding or prediction at each moment in the sequence.

HOW THE RNN ALGORITHM WORKS STEP BY STEP

Step 8: Application-Specific Use

- Depending on the task the RNN is designed for, the final outputs can be used in various ways.
- For example, in language translation, the outputs could be words in another language that form the translation of the input sentence.

Step 9: Learning and Training

- To make accurate predictions, the RNN needs to learn from examples. It goes through a training process where it adjusts its internal parameters (weights) based on the differences between its predictions and the correct answers.

HOW THE RNN ALGORITHM WORKS STEP BY STEP

Step 10: Repeating for New Data

- Once the RNN is trained, it can process new, unseen sequences and make predictions or generate new sequences based on its learned understanding of the patterns in the data.

In summary, the RNN algorithm works by unrolling over time, processing data at each time step, updating its memory, and making predictions based on the learned patterns in sequential data. This process allows RNNs to handle various tasks, like understanding languages, predicting future values, generating music or text, and much more!

APPLICATIONS OF RNNs

RNNs have found applications in various fields, including:

- **Natural Language Processing (NLP):** For tasks like language translation, sentiment analysis, text generation, and speech recognition.
- **Time Series Prediction:** For forecasting stock prices, weather conditions, or any time-dependent data.
- **Music Generation:** For creating music with a sense of coherence and continuity.
- **Video Analysis:** For action recognition and understanding dynamic scenes.

**THANK YOU FOR
LISTENING!**